

CS542200 Parallel Programming

Homework 4: Blocked All-Pairs Shortest Path

104062703 曾若淳

Implementation

對於一個有 N 個點的圖，我使用 blocked Floyd-Warshall來算出所有path的最短路徑。基本上切data的方法就是讓cuda的block大小就等於 blocked Floyd-Warshall演算法中的block大小。

Single-GPU: Cuda

需要計算的 $N \times N$ 個路徑，第一維和第二維各會被分成 $round = \frac{N}{B}$ 個block，演算法在每個iteration 必須按照順序計算三個階段，但每個階段內不同block間可以平行計算。我讓每個block內每個thread計算 B 個

$dist[a][b] = \min(dist[a][b], dist[a][k] + dist[k][b])$ 的最短路徑的更新。

假如直接從global memory存取的話，會需要存取 $3B$ 個位置，整個block會同時存取 $3B^3$ 次，很慢而且不同thread可能會存取重複的位置發生衝突。所以我使用了shared memory，事先將這個block的所有位置 $dist[a][b]$ 以及所有 $dist[a][k]$, $dist[k][b]$ 這三個block的資料放到shared memory裡，這部份可以平行做，讓每個thread存取 $O(1 + 2B)$ 個位置即可。

當這個block所有會用到的資料都放到shared memory後，就可以非常快速的計算最短路徑，算完每個thread在平行的將它負責的那個最短路徑 $dist[a][b]$ 放回global memory裡。

Multi-GPU: OpenMP

有 n 張GPU，我會開 n 個OpenMP thread，每個控制一張GPU，而每個GPU負責計算 $\frac{round}{n} \times round$ 個block，GPU內計算方式和上面的single-GPU版本相同，差別在於每個iteration的三個階段間，必須將GPU內的資料做同步。我實做資料同步的方法為，在host上pinned一塊memory，GPU會先將自己更新的部份放到host上對應的位置，這部份使用openmp的barrier來確保每張GPU都已經放到對應位置，之後GPU才將其他GPU的更新從host拿回device上。

Multi-GPU: MPI

基本上和OpenMP實做非常相似，有 n 張GPU，我會開的 n 個MPI process。每個iteration內每個階段間平行計算路徑的部份也跟single-GPU版本相同，而階段間的資料同步，我使用MPI_Barrier來確保GPU都已經將負責部份的更新到host對應的memory，然後GPU才會load其他GPU負責的部份。

Performance Analysis

實驗環境為課程提供的GPU cluster，有兩張Tesla K20的node。

Profiling Results

我使用nvvp，以下為三個版本在problem size=500, block大小=8,16,32的情況：

Single-GPU

B=8

```
cal(int, int, int**, int, int, int, int, int)
```

Maximum instruction execution count in assembly: 12021

Average instruction execution count in assembly: 3476

Instructions executed for the kernel: 1147254

Thread instructions executed for the kernel: 36712128

Non-predicated thread instructions executed for the kernel: 32220672

Warp non-predicated execution efficiency of the kernel: 87.8%

Warp execution efficiency of the kernel: 100.0%

L1/Shared Memory

Local Loads	0	0 B/s
Local Stores	0	0 B/s
Shared Loads	118563	111.447 GB/s
Shared Stores	112740	105.973 GB/s
Global Loads	131220	15.761 GB/s
Global Stores	11664	1.37 GB/s
Atomic	0	0 B/s
L1/Shared Total	374187	234.552 GB/s

L2 Cache

L1 Reads	134136	15.761 GB/s
L1 Writes	11664	1.37 GB/s

B=16

```
cal(int, int, int**, int, int, int, int, int)
```

Maximum instruction execution count in assembly: 34160

Average instruction execution count in assembly: 7249

Instructions executed for the kernel: 2392190

Thread instructions executed for the kernel: 76547040

Non-predicated thread instructions executed for the kernel: 66118192

Warp non-predicated execution efficiency of the kernel: 86.4%

Warp execution efficiency of the kernel: 100.0%

L1/Shared Memory

Local Loads	0	0 B/s
Local Stores	0	0 B/s
Shared Loads	551725	458.779 GB/s
Shared Stores	514408	427.749 GB/s
Global Loads	556920	59.022 GB/s
Global Stores	29120	3.027 GB/s
Atomic	0	0 B/s
L1/Shared Total	1652173	948.577 GB/s

L2 Cache

L1 Reads	567840	59.022 GB/s
L1 Writes	29120	3.027 GB/s

B=32

```
cal(int, int, int**, int, int, int, int, int)
```

Maximum instruction execution count in assembly: 45471

Average instruction execution count in assembly: 10004

Instructions executed for the kernel: 3301326

Thread instructions executed for the kernel: 105639080

Non-predicated thread instructions executed for the kernel: 90933588

Warp non-predicated execution efficiency of the kernel: 86.1%

Warp execution efficiency of the kernel: 100.0%

Local Loads	0	0 B/s
Local Stores	0	0 B/s
Shared Loads	1728929	471.052 GB/s
Shared Stores	1474683	401.782 GB/s
Global Loads	1528640	52.616 GB/s
Global Stores	43520	1.482 GB/s
Atomic	0	0 B/s
L1/Shared Total	4775772	926.932 GB/s

L2 Cache

L1 Reads	1544960	52.616 GB/s
L1 Writes	43520	1.482 GB/s

Multi-GPU: OpenMP

B=8

```
cal(int, int, int, int**, int, int, int, int, int)
```

Maximum instruction execution count in assembly: 9643

Average instruction execution count in assembly: 2765

Instructions executed for the kernel: 909774

Thread instructions executed for the kernel: 22864128

Non-predicated thread instructions executed for the kernel: 20021376

Warp non-predicated execution efficiency of the kernel: 68.8%

Warp execution efficiency of the kernel: 78.5%

L1/Shared Memory

Local Loads	0	0 B/s
Local Stores	0	0 B/s
Shared Loads	101424	103.93 GB/s
Shared Stores	101382	103.887 GB/s
Global Loads	124416	16.822 GB/s
Global Stores	9216	1.18 GB/s
Atomic	0	0 B/s
L1/Shared Total	336438	225.819 GB/s

L2 Cache

L1 Reads	131328	16.822 GB/s
L1 Writes	9216	1.18 GB/s

B=16

cal(int, int, int, int**, int, int, int, int, int)
--

Maximum instruction execution count in assembly: 4096

Average instruction execution count in assembly: 429

Instructions executed for the kernel: 141396

Thread instructions executed for the kernel: 4379568

Non-predicated thread instructions executed for the kernel: 3983984

Warp non-predicated execution efficiency of the kernel: 88.1%

Warp execution efficiency of the kernel: 96.8%

L1/Shared Memory

Local Loads	0	0 B/s
Local Stores	0	0 B/s
Shared Loads	34285	175.88 GB/s
Shared Stores	26463	135.754 GB/s
Global Loads	21908	15.154 GB/s
Global Stores	1952	1.252 GB/s
Atomic	0	0 B/s
L1/Shared Total	84608	328.04 GB/s

L2 Cache

L1 Reads	23632	15.154 GB/s
L1 Writes	2036	1.306 GB/s

B=32

cal(int, int, int, int**, int, int, int, int, int)
--

Maximum instruction execution count in assembly: 95718

Average instruction execution count in assembly: 18134

Instructions executed for the kernel: 5966161

Thread instructions executed for the kernel: 156113392

Non-predicated thread instructions executed for the kernel: 133974776

Warp non-predicated execution efficiency of the kernel: 70.2%

Warp execution efficiency of the kernel: 81.8%

Local Loads	0	0 B/s
Local Stores	0	107.403 MB/s
Shared Loads	3488157	532.157 GB/s
Shared Stores	3028032	461.96 GB/s
Global Loads	3170816	66.268 GB/s
Global Stores	90112	1.718 GB/s
Atomic	0	0 B/s
L1/Shared Total	9777117	1,062.21 GB/s


L2 Cache

L1 Reads	3486112	66.481 GB/s
L1 Writes	95766	1.826 GB/s


Weak Scalability

single-GPU版本，時間幾乎全花在計算上，而在不同problem size=node數下，可以看到使用大小為16的block表現都最好。


(I/O的部份，音讀寫檔案的時間是固定的，故不列入比較)



Multi-GPU的OpenMP版本，可以觀察到計算時間，由於使用兩張GPU大致上時間減半，但是花在溝通時間、複製資料的時間變長，整體時間主要由複製資料的部份dominate。而block大小對於執行時間的影響，仍可看到在大小為16時能計算的最快，但是由於複製資料跟block大小成反比，所以在block最大=32時，整體執行時間最低。




Multi-GPU的MPI版本，由於實做方式和OpenMP版本非常類似，weak scalability趨勢很類似，最好的block大小也是32。



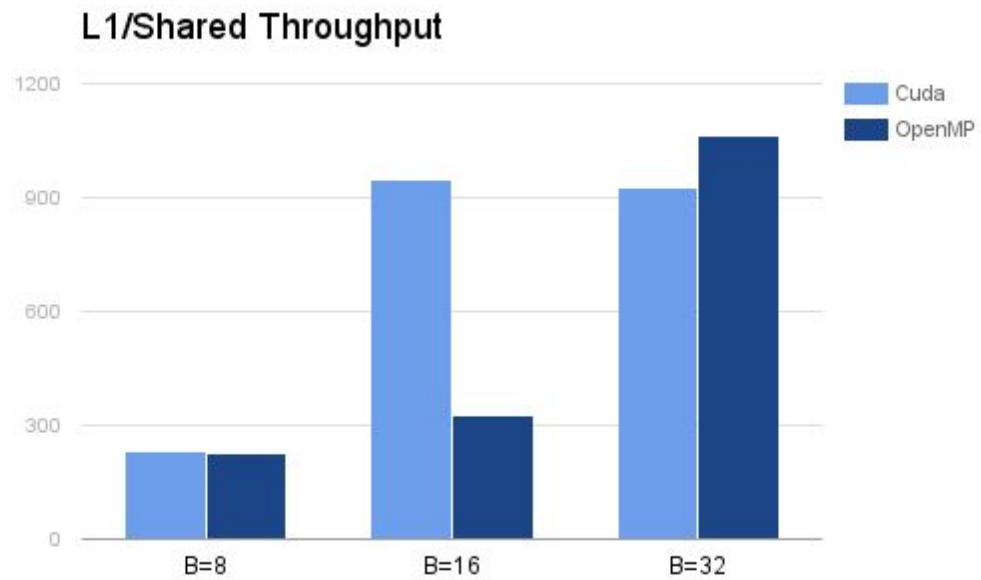
Blocking Factor

GFLOPS

Max Instruction Count in Assembly



Device/Shared Memory Bandwidth




Optimization Techniques

以下實驗皆為problem size=5000個node, block大小為32時的測量結果

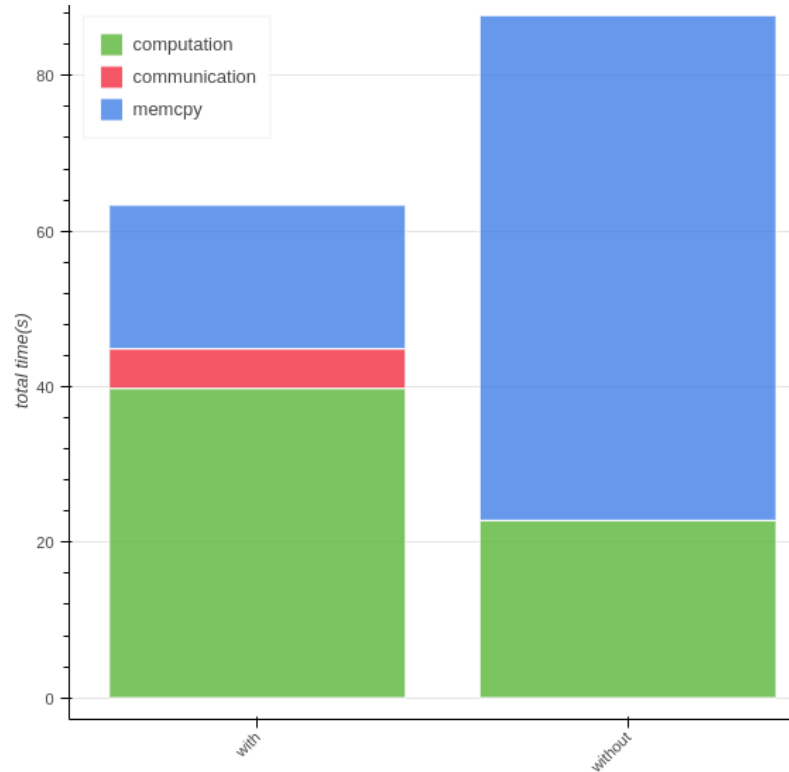
shared memory

在未使用 shared-memory, 每次計算最短路徑都從 global memory存取時所需的時間大約是有使用 shared-memory版本的5倍。



streaming

multi-GPU 版本在不同階段間必須做資料交換，右圖為資料交換使用 streaming 前後的比較圖，沒使用 streaming 平行存取 data 需要的複製資料時間是有用多個 streaming 的 3 倍。



Conclusion

對於計算 All-Pair Shortest Path 問題，使用單一 GPU 可以省去資料同步的時間，但每個階段間卻無法完全利用 data-parallelism 的特性，使用多張 GPU 可以盡量將資料平行處理，但是卻花了更多時間在資料同步上。我想多張 GPU 更適合不需要這麼頻繁的做資料同步，或者能容許延遲同步的問題上（比如 neural network 的 training，晚幾個回合做更新，可能對最終 performance 沒有影響）。